

# Veris AI Candidate Task: Build an AI Agent

---

Your task is to build an AI agent for the financeQA benchmark presented on <https://www.financearena.ai/>. The benchmark currently ranks LLMs in a non-agentic approach, where OpenAI is leading with only 54.1% overall score. The goal is to introduce an agentic approach to improve on this baseline. You can decide to use any tools that you wish including a RAG system to improve on the benchmark evaluation dataset. The task's main focus is on your thought process and implementation rather than the actual score you get on the benchmark.

## Requirements

- Language: Python or TypeScript.
- Frameworks & Models: Any LLM of your choice (e.g., OpenAI, open-source LLMs). You may use commercial APIs or open-source models. You may also choose to use an agentic framework of your choice or not.
- Tools: You can choose any API, MCP, or existing tool library/framework. You can decide to use RAG as a tool.
- You are free to use any coding assistant or environment for the project.

## Deliverables

1. Working code for your agent
2. Documentation explaining:
  - a. An agent card explaining what the agent does
  - b. Your overall design approach
  - c. Key trade-offs and decisions
  - d. Evaluation over the benchmark dataset
  - e. Suggestions for future improvements
3. (Optional) Containerized deployment (e.g., Docker) or host on a cloud environment

## Evaluation Criteria

- Code quality & organization (clarity, modularity, best practices)
- Functionality (agent should work reliably and meet the task objective)
- Documentation quality (clear explanation of design and reasoning)
- (Bonus but not required) Benchmark score and improvement over the current non-agentic baselines

## Expected Effort

This task is designed to take 4 to 8 hours of work. We expect a submission within a week.

## Submission

Email us a GitHub repository link or a zip file containing:

- All source code
- Any configuration or environment files (e.g., requirements.txt or package.json)
- Documentation
- (Optional) Deployment files (Dockerfile, scripts, etc.)

If you have any questions, please reach out to [hello@veris.ai](mailto:hello@veris.ai)